

```
In [1]: print("Hello World!")
```

Hello World!

```
In [2]: # 1.a
import pandas as pd
spreadsheet = pd.read_csv('C:/Users/colto/Downloads/auto.csv')
```

```
In [3]: # 1.b
spreadsheet.head()
```

```
Out[3]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	name
0	18.0	8	307.0	130	3504	12.0	70.0	1	chevrolet chevelle malibu
1	15.0	8	350.0	165	3693	11.5	70.0	1	buick skylark 320
2	18.0	8	318.0	150	3436	11.0	70.0	1	plymouth satellite
3	16.0	8	304.0	150	3433	12.0	70.0	1	amc rebel sst
4	17.0	8	302.0	140	3449	NaN	70.0	1	ford torino

```
In [4]: # 1.c
print("Rows: ", len(spreadsheet))
print("Columns: ", len(spreadsheet.columns))
```

Rows: 392  
Columns: 9

```
In [5]: # 2.a
spreadsheet['mpg'].describe()
spreadsheet['weight'].describe()
spreadsheet['year'].describe()
```

```
Out[5]: count    390.000000
mean      76.010256
std       3.668093
min       70.000000
25%      73.000000
50%      76.000000
75%      79.000000
max       82.000000
Name: year, dtype: float64
```

```
In [6]: # 2.b Range
ss1 = spreadsheet.iloc[:, 0:8]
print(ss1.max() - ss1.min())
```

mpg 37.6  
cylinders 5.0  
displacement 387.0  
horsepower 184.0  
weight 3527.0  
acceleration 16.8  
year 12.0  
origin 2.0  
dtype: float64

```
In [7]: # 2.b Range
print(ss1.max() - ss1.min())
```

mpg 37.6  
cylinders 5.0  
displacement 387.0  
horsepower 184.0  
weight 3527.0  
acceleration 16.8  
year 12.0  
origin 2.0  
dtype: float64

```
In [8]: # 3.a
spreadsheet.dtypes
```

```
Out[8]: mpg                float64
cylinders                int64
displacement            float64
horsepower              int64
weight                  int64
acceleration            float64
year                    float64
origin                  int64
name                     object
dtype: object
```

```
In [9]: # 3.b
spreadsheet['cylinders'] = spreadsheet['cylinders'].astype('category').cat.codes
```

```
In [10]: # 3.c
spreadsheet['origin'] = spreadsheet['origin'].astype('category')
```

```
In [11]: # 3.d
spreadsheet.dtypes[['cylinders', 'origin']]
```

```
Out[11]: cylinders        int8
origin          category
dtype: object
```

```
In [12]: # 4.a
spreadsheet = spreadsheet.dropna()
```

```
In [13]: # 4.b
print("Rows: ", len(spreadsheet)) # Originally 392
print("Columns: ", len(spreadsheet.columns)) # Originally 9

Rows: 389
Columns: 9
```

```
In [14]: # 5.a
avg_mpg = spreadsheet['mpg'].mean()
spreadsheet['mpg_high'] = 0
spreadsheet.loc[spreadsheet['mpg'] > avg_mpg, 'mpg_high'] = 1
```

```
In [15]: # 5.b
spreadsheet = spreadsheet.drop(columns=['mpg', 'name'])
```


```
In [16]: # 5.c
spreadsheet.head()
```

```
Out[16]:
```

	cylinders	displacement	horsepower	weight	acceleration	year	origin	mpg_high
0	4	307.0	130	3504	12.0	70.0	1	0
1	4	350.0	165	3693	11.5	70.0	1	0
2	4	318.0	150	3436	11.0	70.0	1	0
3	4	304.0	150	3433	12.0	70.0	1	0
6	4	454.0	220	4354	9.0	70.0	1	0


```
In [17]: # 6.a
import seaborn as sns
sns.catplot(spreadsheet['mpg_high'])
```

```
Out[17]: <seaborn.axisgrid.FacetGrid at 0x278173bee50>
```




```
In [18]: # 6.b
sns.relplot(x='horsepower', y='weight', hue='mpg_high', data=spreadsheet)
```

```
Out[18]: <seaborn.axisgrid.FacetGrid at 0x278174fd3d0>
```



```
In [19]: # 6.c
sns.boxplot(x='horsepower', y='weight', data=spreadsheet)
```

```
Out[19]: <Axes: xlabel='horsepower', ylabel='weight'>
```



```
In [20]: # 6.d
# The graph from 6.a shows that only the values 0 and 1 exist inside the mpg_high category.
# The graph from 6.b shows that observations with a mpg_high value of 1 tend to have low horsepower and weight compared to observations with a mpg_high value of 0.
# The graph from 6.c shows that outliers tend to have high weight and low horsepower rather than low weight and high horsepower.
```

```
In [21]: # 7.a
split = 0.2
```

```
In [22]: # 7.b
seed = 1234
```

```
In [23]: # 7.c
from sklearn.model_selection import train_test_split

data = spreadsheet
x = data.iloc[:, 0:7] # x = feature matrix
y = data['mpg_high'] # y = target variable
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=split, random_state=seed)
```

```
In [24]: # 7.d
print("x_train shape:", x_train.shape)
print("x_test shape:", x_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)

x_train shape: (311, 7)
x_test shape: (78, 7)
y_train shape: (311,)
y_test shape: (78,)
```

```
In [25]: # 8.a
from sklearn.linear_model import LogisticRegression

model = LogisticRegression(solver='lbfgs', max_iter=1000)
model.fit(x_train, y_train)
y_pred = model.predict(x_test)
```

```
In [26]: # 8.b
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1_score = f1_score(y_test, y_pred)

print("Accuracy: {:.2f}".format(accuracy))
print("Precision: {:.2f}".format(precision))
print("Recall: {:.2f}".format(recall))
print("F1 Score: {:.2f}".format(f1_score))
```

Accuracy: 0.90  
Precision: 0.78  
Recall: 1.00  
F1 Score: 0.88

```
In [27]: # 8.c
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.84	0.91	50
1	0.78	1.00	0.88	28
accuracy			0.90	78
macro avg	0.89	0.92	0.89	78
weighted avg	0.92	0.90	0.90	78

```
In [28]: # 9.a
from sklearn.tree import DecisionTreeClassifier

dt_model = DecisionTreeClassifier()
dt_model.fit(x_train, y_train)
```

```
Out[28]: ▾ DecisionTreeClassifier
DecisionTreeClassifier()
```


```
In [29]: # 9.b
y_pred = dt_model.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy: {:.2f}".format(accuracy))
```

Accuracy: 0.90

```
In [30]: # 9.c
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.86	0.91	50
1	0.79	0.96	0.87	28
accuracy			0.90	78
macro avg	0.89	0.91	0.89	78
weighted avg	0.91	0.90	0.90	78

```
In [31]: # 9.d
from sklearn import tree
import matplotlib.pyplot as plt
tree.plot_tree(dt_model)
plt.show()
```



```
In [32]: # 10.a
from sklearn.neural_network import MLPRegressor
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
scaler1 = StandardScaler()
scaler2 = StandardScaler()
scaler.fit(x_train)
scaler1.fit(x_test)

x_train_scaled = scaler.transform(x_train)
x_test_scaled = scaler1.transform(x_test)

regr = MLPRegressor(hidden_layer_sizes=(6,3), max_iter=1000, random_state=seed)
regr.fit(x_train_scaled, y_train)
regr.fit(x_test_scaled, y_test)
```

```
Out[32]: MLPRegressor
MLPRegressor(hidden_layer_sizes=(6, 3), max_iter=1000, random_state=1234)
```

```
In [33]: # 10.b
from sklearn.metrics import mean_squared_error, r2_score

y_pred_regr = regr.predict(x_test_scaled) # make predictions

print('mse =', mean_squared_error(y_test, y_pred_regr))
print('correlation =', r2_score(y_test, y_pred_regr))

mse = 0.05527247016944208
correlation = 0.7598016367779388
```

```
In [34]: # 10.c

scaler = StandardScaler()
scaler1 = StandardScaler()
scaler.fit(x_train)
scaler1.fit(x_test)

x_train_scaled = scaler.transform(x_train)
x_test_scaled = scaler1.transform(x_test)

regr = MLPRegressor(hidden_layer_sizes=(12,6), max_iter=1000, random_state=seed, solver='adam') # different topology, different
regr.fit(x_train_scaled, y_train)
regr.fit(x_test_scaled, y_test)
```

```
Out[34]: MLPRegressor
MLPRegressor(hidden_layer_sizes=(12, 6), max_iter=1000, random_state=1234)
```

```
In [35]: # 10.d
y_pred_regr = regr.predict(x_test_scaled) # make predictions

print('mse =', mean_squared_error(y_test, y_pred_regr))
print('correlation =', r2_score(y_test, y_pred_regr))

mse = 0.04772849165185324
correlation = 0.7925856119929463
```

```
In [36]: # 10.e
# After playing around with different settings, it was difficult to get a healthier mse than the first model.
# But if I increase the size of the topology while retaining the same 2:1 node ratio, I can get an MSE lower than model 1.
# It makes sense that both models perform relatively the same because their topology & settings are similar.
# However, as I increase the number of nodes, I risk overfitting the data.
```

```
In [37]: # 11.a
# Neural networks took the longest to execute. Logistic regression executed just about as fast as the decision tree.
# However, decision trees looks like it reports better metrics than the decision tree. (Investigated in 11.b)
```

```
In [38]: # 11.b
# Logistic Regression
#
#           precision    recall  f1-score   support
#
#      0         1.00        0.84        0.91         50
#      1         0.78        1.00        0.88         28
#
#   accuracy                0.90         78
#  macro avg         0.89        0.92        0.89         78
#weighted avg         0.92        0.90        0.90         78
# Decision Tree
#
#           precision    recall  f1-score   support
#
#      0         0.94        0.92        0.93         50
#      1         0.86        0.89        0.88         28
#
#   accuracy                0.91         78
#  macro avg         0.90        0.91        0.90         78
#weighted avg         0.91        0.91        0.91         78
#
# It's difficult to say which performed better because their numbers are similar.
# But it looks like the decision tree reports better metrics for f1-score while precision and recall are relatively the same.
```

```
In [39]: # 11.c
# In the areas where decision trees might have outperformed logistic regression, it can be partially credited to it's robustness.
# Alternatively, there may be some interaction between features which decision trees are great at capturing.
# On the other hand, there might not have been much interaction so logistic regression was able to keep up with decision trees.
```

```
In [40]: # 11.d
# sklearn was nice to pick up since I'm familiar with python, but my kernel had crashed a couple times so far.
# It's a little troublesome since I wouldn't immediately realize that it's stuck or crashed.
# I still prefer sklearn, I especially like having my development environment in my browser so it's easier to look up information.
```